


MIXED (REALISTIC) WORKLOAD

Aggregate metrics comparison of between AI-augmented Linux-like baseline scheduler and the Linux-like baseline scheduler:

Metric	Linux-like Baseline Scheduler	AI-augmented Linux-like Scheduler
Simulated ticks	28410	28410
Average Turnaround Time	13337.8394	9749.6503
Median Turnaround Time	9009.0	3893.0
Average Response Time	13308.3015	9718.0676
95th Percentile Response	26984.5	26984.5
Fairness (Jain Index)	0.003796679	0.003796679
CPU Utilization (%)	99.876808	99.876808
Context Switches	103261	100288
Tasks Total	1227	1227
Tasks Completed	1227	1227
Throughput (tasks/unit time)	0.043187	0.043187

PLOT:



Results Summary —

On a realistic mixed workload of 1,227 tasks (file editing, browsing, video, downloads, background jobs), the AI-augmented scheduler showed strong performance improvements:

- **Average Turnaround Time improved by 27% (13,338 → 9,750 ticks).**
- **Median Turnaround Time dropped by 57% (9,009 → 3,893 ticks), highlighting smoother handling of everyday interactive tasks.**
- **Average Response Time reduced by 27% (13,308 → 9,718 ticks).**

- 95th percentile response time stayed constant, ensuring worst-case stability.
- Fairness Index and CPU Utilization matched Linux baseline (no sacrifice in fairness or efficiency).
- **Context Switches decreased by ~2.9%, lowering scheduling overhead.**

Interpretation: In real-world multitasking conditions, AVIOS consistently improves responsiveness and turnaround while keeping system fairness and efficiency intact.